

Suleyman Kiani

Applied AI Engineer · Agentic Systems, Backend & Formal Verification

Burlington, ON | +1 (289) 788-8260 | suley.kiani@outlook.com | linkedin.com/in/suleyman-kiani | github.com/kianis4 | suleyman.io

SUMMARY

Applied AI engineer who ships agents with their evals attached. A year inside a regulated commercial lender taking systems from business discovery through IT intake and security review to corporate handoff, alongside a founder's multi-tenant payments platform that a real business runs on. MEng in Computing & Software at McMaster, where my project formally verifies that an LLM pipeline cannot release an uncited claim or leak a confidential one.

TECHNICAL SKILLS

Languages: Python, TypeScript, JavaScript, Java, SQL, Swift, C

AI & Agents: LLM APIs (Anthropic, OpenAI), agent orchestration, LangGraph, structured outputs, RAG, vector search (LanceDB, MongoDB Atlas), BGE-M3 embeddings, Ollama, evaluation harnesses, LLM-as-judge

Backend & Data: REST/JSON APIs, microservices, event-driven architecture, RabbitMQ, Redis Streams, WebSockets, Node.js, FastAPI, Prisma, PostgreSQL, query tuning and indexing, schema migrations, Google Workspace APIs

Cloud & DevOps: AWS Lambda, Vercel, Docker, Linux, systemd, kernel-namespaces sandboxing, Jenkins, GitHub Actions, CI/CD, Gitleaks secret scanning

Testing & Formal Methods: TDD, pytest, Vitest, Playwright, JUnit, mutation testing, TLA+ model checking, Z3/SMT2

Domain: Regulated financial services, equipment and lease finance, credit underwriting submissions, financial statement analysis, NPV and IRR math, collections risk scoring

EXPERIENCE

Applied AI Engineer

Starting Oct. 2026

LaunchGood

Remote

- Senior (P3) hire into the pod that owns AI initiatives at a global crowdfunding platform: production agents and automations across operations, trust and safety, support and finance, each gated by **evaluation suites in CI** (deterministic checks, LLM-as-judge, human review). The hiring bar was a deployed, evaluated prototype: Zakat-Eligibility Triage, below.

Associate Account Manager (Software Engineering scope)

Sept. 2025 – Oct. 2026

Mitsubishi HC Capital Canada

Burlington, ON

- Built one bilingual **TypeScript** application to replace **13 Excel pricing calculators** across 14 manufacturers and 9 pricing engines; rate configs are versioned **Postgres** rows with an audit log, so repricing never needs a code deploy.
- Carried it through **enterprise SDLC**: formal IT intake (a nine-document security-scoped submission), a leadership review, an IT pre-review, then a handoff of source, calculation spec and **16 written rulings** (capital and FMV leases, loans, residuals, cost-of-funds ladders) to corporate IT, which now owns it and is deploying it from that documentation.
- Validated its finance core (payment, rate and IRR functions) against the booking system of record: **45 funded deals** reproduced to within half a cent given each deal's booked rate and structure, so the arithmetic is proven on real flows, not its own output.
- Gated it behind a **1,370-test** suite with mutation-tested assertions, funded-deal fixtures kept server-only and expected-fail tests holding known gaps open; it caught **7 of 9 engines** silently quoting wrong terms.
- Co-designed, with two colleagues, a collections risk scorecard from six years of delinquency history: five bands separate average default probability **0.647 to 0.017** on the scored live book, a 38x spread with no reversals; model selection used customer-grouped cross-validation.
- Approved for launch and announced org-wide by the SVP in Sept. 2026 as a risk-based prioritization tool after it re-scored a 1,748-contract book, ranking 1,023 contracts carrying **\$66.4M** of net investment; collections managers now work from its **daily list**.
- Built the **AI document-validation service**: LLM extraction checked against a documented rules checklist, behind a deterministic intake-and-retry state machine, returning a READY / NOT READY verdict on a deal package at a **median of 4.9 minutes** over 56 runs; advisory only, zero incremental spend on data already held.

Software Engineer & Founder

Jan. 2024 – Present

SKomp Studio

Remote

- Built **Incite**, a white-label multi-tenant booking-and-payments platform on a **74-model Prisma schema** and **196 REST API routes**: scheduling, waitlists, Square payments, memberships, ticketing and per-tenant theming.
- Its flagship tenant runs its entire operation on it: **\$45K+ CAD** processed since launch across **640 registered users**, 1,380 confirmed bookings and 105 active memberships, all self-serve with no parallel manual system.
- Money is the part that cannot be wrong: **idempotent payments**, webhook signature verification and **tenant isolation** enforced in a tenant-scoped data-access layer, guarded by a build-failing test for unscoped models.
- Run it to enterprise standards: **6,000+ automated tests** gating every merge, Playwright suites as required status checks, **Gitleaks** secret scanning in CI, and branch-protected pull requests with a preview deploy per change.
- Mentored **100+ learners** through algorithms, systems design and interview preparation, running every engagement end to end since 2024: acquisition, pricing, curriculum design, live sessions and follow-up review.

Junior Web Developer

May 2021 – Jun. 2022

Giftcash Inc.

Remote

- As one of three engineers in a 30 to 40 person development organization, migrated a legacy Django monolith to Node.js on **AWS Lambda** for per-request billing, tuned PostgreSQL indexes to kill the sequential scans behind balance lookup and stood up Jenkins and GitHub Actions CI/CD.

EDUCATION

McMaster University

Hamilton, ON

MEng, Computing & Software (Distributed Systems, Formal Methods & Applied ML) Sept. 2025 – 2027 (expected)

- Mac Study Companion**, co-supervised by Dr. Richard Paige and Dr. William Farmer and judged complete by both in Sept. 2026: a lecture-to-study-notes pipeline in which no claim reaches a student until it cites its exact source. A Z3/SMT encoder emits a four-conjunct verification condition (completeness, support gating, non-interference over a confidentiality lattice, recall floor); the gate refuses if any conjunct fails.
- Model-checked the pipeline in TLA+ (4 properties over 2,647 states at the shipped configuration), with a deliberately broken mutant spec that must fail in CI so a green check is never vacuous. On a 10-lecture corpus the gate took citation precision from 0.604 to **0.960**; recall 1.000 and leak rate **0.000** hold by construction, since the gate's conjuncts are the predicates being measured.
- Delivered as **9 hexagonal microservices** over Redis Streams behind 2,489 backend tests, with a 32-page project report. Earlier built **PodcastHub**, six event-driven services over a RabbitMQ topic exchange. **A+** in Simple Type Theory, **A** in Microservice Architectures.

McMaster University

Hamilton, ON

BASc (Honours), Computer Science

Sept. 2018 – Nov. 2024

PROJECTS

Zakat-Eligibility Triage · evaluated LLM triage prototype | *Python, structured outputs, LLM-as-judge* 2026

- An eligibility-triage agent for crowdfunding campaigns, scored by its own harness: 18 hand-labelled campaigns, an LLM judge blind to the labels and a human call recorded on every case, so no verdict ships unreviewed; **10 decision records** carry the tradeoffs.
- Scored **97.2%** category agreement (140 of 144) and **83.3%** escalation agreement (15 of 18), clearing gates of 80% and 75%, behind 442 unit tests; five of the 18 cases are labelled ambiguous on purpose, so a perfect score would be the red flag.

Personal Agentic OS · distributed automation platform | *Python, TypeScript, Docker, Linux, systemd* 2026

- Built and run a self-hosted automation platform: **60+ scheduled services** and event-triggered workers on Linux; the workers that take untrusted input run in kernel-namespaces sandboxes behind a fail-closed egress allowlist.
- Work that dies mid-flight is recovered rather than lost: a **dead-man switch** watches the health monitor itself, and a sweeper re-drives stranded jobs up to three times before parking them for review instead of retrying forever.
- A local retrieval stack (BGE-M3 embeddings, gemma3n on an RTX 3080, LanceDB) returns cited answers tuned on a hand-built **40-question Recall@10** gold set; no document, query or answer leaves the machine.

Mike Ross AI · legal retrieval assistant (Delta Hacks 12) | *Python, FastAPI, MongoDB Atlas, LangGraph* 2026

- Answers Canadian-law questions with quote-grounded citations: a LangGraph state machine runs jurisdiction-aware retrieval over a **16,000-chunk** corpus spanning seven statutes, keeps multi-turn memory and drafts send-ready notices to PDF; built in a 24-hour hackathon.

SKomp Forge · agentic iOS training coach | *Swift 6, SwiftUI, SwiftData, HealthKit*

2025 – Present

- Rebuilt my training app into an **agentic coach** across 11 Swift 6 modules, shipping to TestFlight: 36 tool schemas pinned in both Swift and Python and diff-tested for drift, every write a typed proposal the lifter confirms by tap, and full-duplex voice holding **P90 313ms** text-to-first-audio under a 1,500ms budget.